A10

A10 AI-Inference Guardian

A strategic infrastructure efficiency layer that reclaims wasted GPU capacity to eliminate the network tax on your enterprise AI stack.

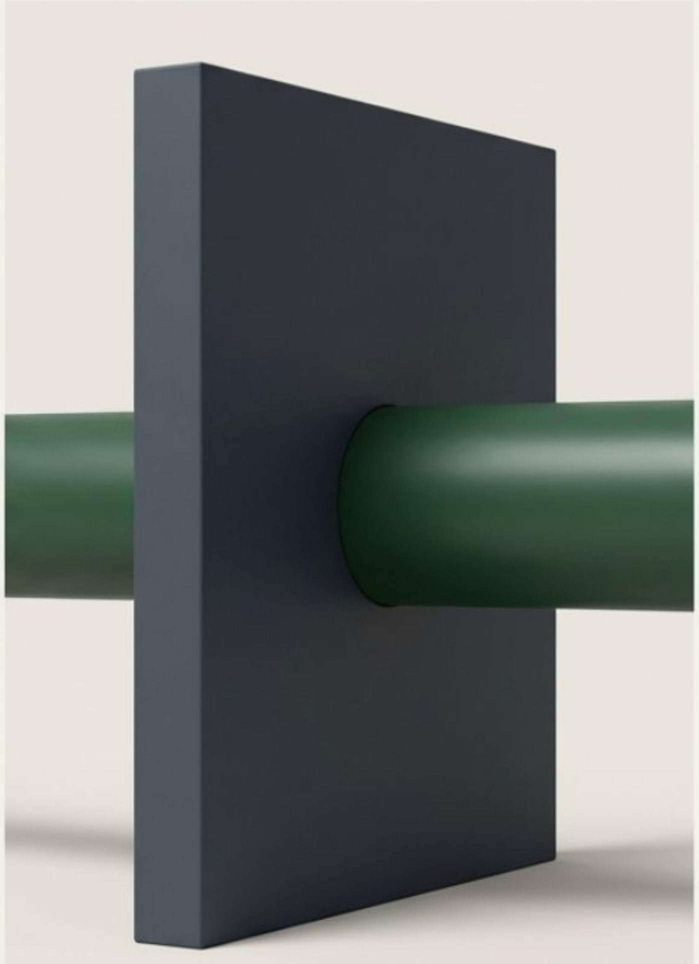

Hitting the GPU wall

The Economic Reality

Enterprise compute demand is rapidly exceeding both budget and available data center power constraints.

The Network Tax

Massive capital investments in H100 and B200 clusters are failing to deliver optimal capacity because the underlying infrastructure is choked by traditional networking overhead.

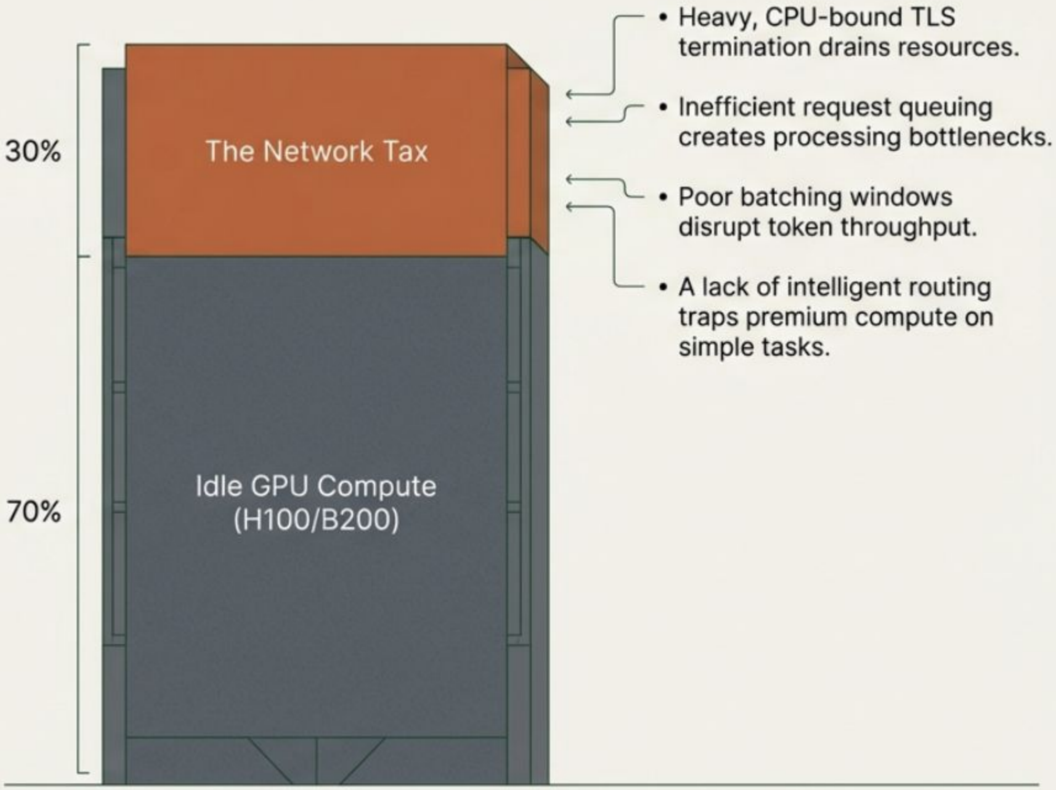

Reclaiming backend capacity at the network edge

“Up to 30% effective backend GPU capacity reclaimed.”

By sitting at the front door of the AI data center.

Efficiency

Lowers the absolute cost per inference by deferring multi-million dollar CapEx expansions and recovering wasted GPU cycles.

Experience

Drives user adoption through faster Time to First Token (TTFT) and stable output generation, without degrading legacy web application performance.

Governance

Centralizes enterprise visibility, enforcing token usage quotas and budget upper caps to eliminate unexpected bill shock.

Moving beyond legacy Layer 7 load balancing

The Limits of the Old Model

Traditional cloud-native gateways (like NGINX or Envoy) were designed for stateless, small-packet web traffic. LLM traffic breaks this model with massive prompt payloads and token-by-token response streaming.

The Hardware Advantage

A10's proprietary FPGA hardware natively handles resource-intensive TLS 1.3 decryption. This offloads the encryption burden entirely, freeing up expensive compute cycles for actual AI inference.

Layer 8 Semantic Routing

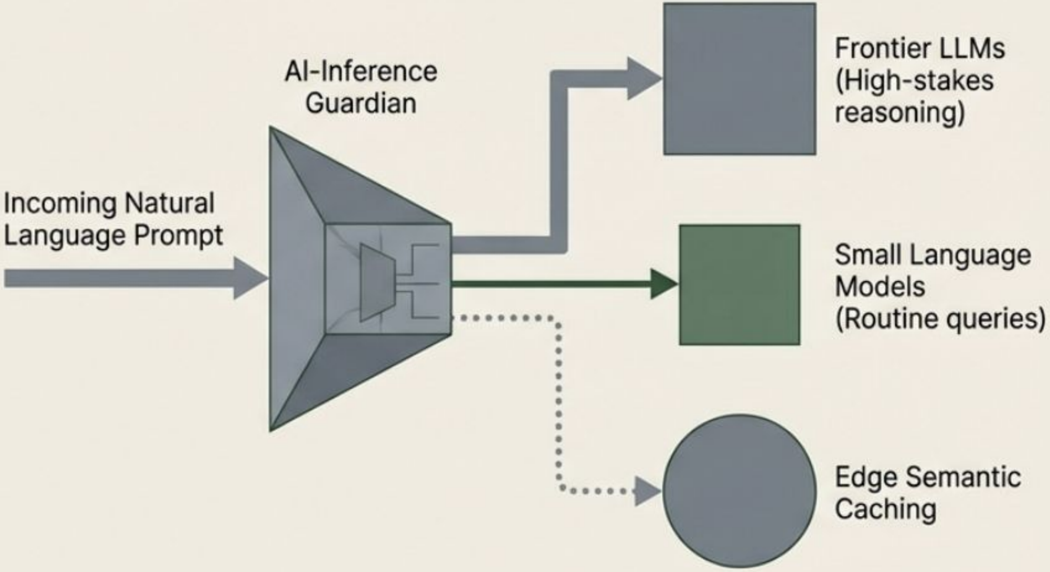

Instead of blindly routing packets, the Guardian analyzes actual intent and routes based on real-time GPU metrics.

Inline defense for semantic risks

The Gap in Traditional Security: Standard Web Application Firewalls (WAFs) rely on strict pattern matching. They are entirely blind to the semantic vulnerabilities introduced by LLMs, such as prompt injection and contextual data leakage.

The AI Firewall

Enforces prompt tokenization, real-time PII redaction, jailbreak pattern detection, and audit logging directly at the network edge.

Healthcare: Redacts 18 distinct PII/PHI identifiers in real-time to maintain strict HIPAA 'Safe Harbor' status.

LEDGER

Fintech: Enforces SEC/FINRA-level audit compllance on all AI-generated communications.

Legal: Preserves attorney-client privilege by ensuring highly sensitive data remains strictly within a secure, closed-system network boundary.

Building the product through structured customer signals

Engaging early adopters via Customer Advisory Boards (CABs) and Proof of Concept benchmarks to validate economic fit and integration viability before writing code.

Persona	Signals
• CFO / VP Infra	Excess GPU Cost
• AI Engineer	High Latency TTFT
• Enterprise Architect	Opaque AI Traffic
• CIO / Counsel	Compliance Risk
• CISO	Semantic Risk

Cross-referencing qualitative user interviews with quantitative performance data (cache hit rates, routing distribution, TPS stability) to identify the root cause of customer friction.

Driving engineering spend with objective math

$$\text{Feature Priority Score} = \frac{(\text{Reach} \times \text{Impact} \times \text{Strategic Alignment} \times \text{Revenue Influence})}{\text{Effort}}$$

The Variables

- Reach: Total inference volume share affected.
- Impact: Avoidance of SLA penalties or delta in GPU utilization.
- Alignment: Strategic multiplier for core efficiency, security, or governance.
- Revenue Influence: ARR exposure and renewal timing.

Quantitative Guardrails

Removing guesswork by triggering immediate backend refactoring if telemetry hits specific operational red lines:

Time to First Token (TTFT) Regression: >10%

GPU Utilization Loss: >5%

False Negative Rate: >1%

Cache Hit Rate Deviation: >15%

Aligning product vision with operational execution

Weekly Tactical (Triage Council)

Refines the engineering backlog based on real-time feedback, allocates immediate sprint resources, and conducts root-cause analysis on performance regressions.

Monthly Technical (Operational Signal Review)

Audits system infrastructure health, analyzing TPS stability, GPU utilization trends, and AI Firewall accuracy against newly identified threat vectors.

Quarterly Executive (Strategic Revenue Review)

Evaluates overall engineering output strictly against core business KPIs and market positioning.

Mid-Year Strategic (Mid-Course Correction)

Re-aligns the broader product roadmap based on three specific pivot signals:

1. Market adoption delta
2. Resource velocity check
3. Efficiency thresholds

Funding today's performance and tomorrow's disruption

The 4-Quarter Progression

Q1	Q2	Q3	Q4
Q1 Core Performance	Q2 Strategic Integration	Q3 Efficiency & Governance	Q4 Ecosystem Dominance
Refining Layer 8 routing and latency reduction.	Deploying native support for Claude, Sonnet, and DeepSeek.	Delivering semantic cache optimization and FinOps dashboards.	Expanding custom LLM APIs and enterprise-wide token quota management.

Resource Allocation (70/20/10 Model)

70% Core (0–12 Months)
Maximizing current ADC performance, latency reduction, and reliability.

20% Emerging (12–24 Months)
Scaling into new high-growth verticals with advanced FinOps analytics.

10% Future (24+ Months)
Innovating next-gen infrastructure, including autonomous routing and cross-cloud GPU orchestration.

Measuring success at the enterprise network boundary

The Hard KPIs:

Every engineering hour and architectural decision maps directly back to three fundamental metrics:

1. **Cost-per-inference improvement:** Lowering the barrier to scale enterprise AI.
2. **GPU utilization uplift:** Maximizing the return on existing infrastructure investments.
3. **Security incident reduction:** Operating securely within highly regulated verticals.

A10 Networks provides the vendor-agnostic intelligence and control plane required to securely dominate the Agentic Era.

Why A10 Wins: The Competitive Moat

- **Vs. Legacy ADCs:** Optimized for massive RAG payloads and token-streaming, not just web traffic.
- **Vs. Cloud Gateways:** Provides a vendor-agnostic control plane to eliminate hyperscaler lock-in.
- **The Hardware Edge:** Proprietary FPGA offload eliminates the "Latency Tax" of software-only rivals.
- **Edge Control:** Proactively blocks semantic threats before they reach the expensive GPU cluster.

Scaling from Efficiency to Ecosystem Dominance

Phase 1 (Land):
Reclaim 30% GPU capacity to solve the CFO's immediate ROI crisis.

Phase 2 (Expand):
Activate the AI Firewall to unblock adoption in Healthcare and Fintech.

Phase 3 (Dominate):
Become the "Agentic Era" OS via autonomous routing and cross-cloud orchestration.

The Goal:

Transform AI infrastructure from a cost center into a strategic value driver.

Thank You!

Q & A